

Horse_or_Human

June 21, 2020

```
[1]: import os
import zipfile
```

```
[2]: import tensorflow as tf
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from tensorflow import keras
```

```
[3]: local_zip = 'F:/Courses/COURSERA/Tensorflow in practice/Course1/Week3/
↳horse-or-human.zip'
zip_ref = zipfile.ZipFile(local_zip, 'r')
zip_ref.extractall('/Week3/horse-or-human')
zip_ref.close()
```

```
[4]: # Direactory with training horse pictures
train_horse_dir = os.path.join('/Week3/horse-or-human/horses')
# Directory with training human pictures
train_human_dir = os.path.join('/Week3/horse-or-human/humans')
```

```
[5]: train_horse_names = os.listdir(train_horse_dir)
print(train_horse_names[:10])

train_human_names = os.listdir(train_human_dir)
print(train_human_names[:10])
```

```
['horse01-0.png', 'horse01-1.png', 'horse01-2.png', 'horse01-3.png',
'horse01-4.png', 'horse01-5.png', 'horse01-6.png', 'horse01-7.png',
'horse01-8.png', 'horse01-9.png']
['human01-00.png', 'human01-01.png', 'human01-02.png', 'human01-03.png',
'human01-04.png', 'human01-05.png', 'human01-06.png', 'human01-07.png',
'human01-08.png', 'human01-09.png']
```

```
[6]: print('total training horse images:', len(os.listdir(train_horse_dir)))
print('total training human images:', len(os.listdir(train_human_dir)))
```

```
total training horse images: 500
total training human images: 527
```

```
[7]: %matplotlib inline

import matplotlib.pyplot as plt
import matplotlib.image as mpimg

# Parameters for our graph; we'll output images in a 4x4 configuration
nrows = 4
ncols = 4

# Index for iterating over images
pic_index = 0
```


```
[8]: # Set up matplotlib fig, and size it to fit 4x4 pics
fig = plt.gcf()
fig.set_size_inches(ncols * 4, nrows * 4)

pic_index += 8
next_horse_pix = [os.path.join(train_horse_dir, fname)
                  for fname in train_horse_names[pic_index-8:pic_index]]
next_human_pix = [os.path.join(train_human_dir, fname)
                  for fname in train_human_names[pic_index-8:pic_index]]

for i, img_path in enumerate(next_horse_pix+next_human_pix):
    # Set up subplot; subplot indices start at 1
    sp = plt.subplot(nrows, ncols, i + 1)
    sp.axis('Off') # Don't show axes (or gridlines)

    img = mpimg.imread(img_path)
    plt.imshow(img)

plt.show()
```



```
[22]: from keras import layers
      from keras.models import load_model
      from tensorflow.keras.optimizers import RMSprop
      from tensorflow.keras.preprocessing.image import ImageDataGenerator
      from PIL import Image
      from keras.preprocessing import image
```

```
[10]: model = keras.Sequential([
      tf.keras.layers.Conv2D(16, (3, 3), activation = 'relu', input_shape = (300, 300, 3)),
      tf.keras.layers.MaxPooling2D(2, 2),
      tf.keras.layers.Conv2D(32, (3, 3), activation = 'relu'),
      tf.keras.layers.MaxPooling2D(2, 2),
```

```

tf.keras.layers.Conv2D(64, (3, 3), activation = 'relu'),
tf.keras.layers.MaxPooling2D(2, 2),
tf.keras.layers.Conv2D(64, (3, 3), activation = 'relu'),
tf.keras.layers.MaxPooling2D(2, 2),
tf.keras.layers.Conv2D(64, (3, 3), activation = 'relu'),
tf.keras.layers.MaxPooling2D(2, 2),
tf.keras.layers.Flatten(),
tf.keras.layers.Dense(512, activation = 'relu'),
tf.keras.layers.Dense(1, activation = 'sigmoid')
])

```

```
[11]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 298, 298, 16)	448
max_pooling2d (MaxPooling2D)	(None, 149, 149, 16)	0
conv2d_1 (Conv2D)	(None, 147, 147, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 73, 73, 32)	0
conv2d_2 (Conv2D)	(None, 71, 71, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 35, 35, 64)	0
conv2d_3 (Conv2D)	(None, 33, 33, 64)	36928
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_4 (Conv2D)	(None, 14, 14, 64)	36928
max_pooling2d_4 (MaxPooling2D)	(None, 7, 7, 64)	0
flatten (Flatten)	(None, 3136)	0
dense (Dense)	(None, 512)	1606144
dense_1 (Dense)	(None, 1)	513

Total params: 1,704,097
 Trainable params: 1,704,097
 Non-trainable params: 0

```
[12]: model.compile(loss='binary_crossentropy',  
                    optimizer=RMSprop(lr=0.001),  
                    metrics=['accuracy'])
```

```
[13]: train_datagen = ImageDataGenerator(rescale = 1 / 255)  
train_generator = train_datagen.flow_from_directory(  
    '/Week3/horse-or-human',  
    target_size = (300, 300),  
    batch_size = 128,  
    class_mode = 'binary'  
)
```

Found 1027 images belonging to 2 classes.

```
[15]: history = model.fit(  
    train_generator,  
    steps_per_epoch = 8,  
    epochs = 25,  
    verbose = 1  
)
```

Epoch 1/25

8/8 [=====] - 44s 6s/step - loss: 0.6479 - accuracy: 0.7063

Epoch 2/25

8/8 [=====] - 43s 5s/step - loss: 0.6058 - accuracy: 0.7063

Epoch 3/25

8/8 [=====] - 41s 5s/step - loss: 0.5449 - accuracy: 0.7820

Epoch 4/25

8/8 [=====] - 41s 5s/step - loss: 0.7051 - accuracy: 0.8398

Epoch 5/25

8/8 [=====] - 42s 5s/step - loss: 0.2895 - accuracy: 0.8676

Epoch 6/25

8/8 [=====] - 52s 6s/step - loss: 0.1596 - accuracy: 0.9477

Epoch 7/25

8/8 [=====] - 45s 6s/step - loss: 0.6036 - accuracy: 0.9132

Epoch 8/25

8/8 [=====] - 51s 6s/step - loss: 0.2052 - accuracy: 0.9404

Epoch 9/25

8/8 [=====] - 51s 6s/step - loss: 0.2642 - accuracy: 0.9266

Epoch 10/25
8/8 [=====] - 45s 6s/step - loss: 0.1007 - accuracy: 0.9633

Epoch 11/25
8/8 [=====] - 40s 5s/step - loss: 0.0520 - accuracy: 0.9822

Epoch 12/25
8/8 [=====] - 40s 5s/step - loss: 0.8376 - accuracy: 0.8865

Epoch 13/25
8/8 [=====] - 43s 5s/step - loss: 0.0886 - accuracy: 0.9655

Epoch 14/25
8/8 [=====] - 48s 6s/step - loss: 0.0368 - accuracy: 0.9922

Epoch 15/25
8/8 [=====] - 40s 5s/step - loss: 0.0285 - accuracy: 0.9911

Epoch 16/25
8/8 [=====] - 41s 5s/step - loss: 0.0302 - accuracy: 0.9922

Epoch 17/25
8/8 [=====] - 40s 5s/step - loss: 0.4362 - accuracy: 0.8765

Epoch 18/25
8/8 [=====] - 43s 5s/step - loss: 0.0549 - accuracy: 0.9855

Epoch 19/25
8/8 [=====] - 48s 6s/step - loss: 0.0109 - accuracy: 0.9980

Epoch 20/25
8/8 [=====] - 40s 5s/step - loss: 0.0087 - accuracy: 0.9978

Epoch 21/25
8/8 [=====] - 41s 5s/step - loss: 0.0019 - accuracy: 1.0000

Epoch 22/25
8/8 [=====] - 41s 5s/step - loss: 7.2441e-04 - accuracy: 1.0000

Epoch 23/25
8/8 [=====] - 45s 6s/step - loss: 0.3931 - accuracy: 0.9366

Epoch 24/25
8/8 [=====] - 56s 7s/step - loss: 0.2940 - accuracy: 0.9433

Epoch 25/25
8/8 [=====] - 45s 6s/step - loss: 0.1725 - accuracy: 0.9544

```
[16]: model.save("horse_or_human.h5")
      print("Saved Model to Disk")
```

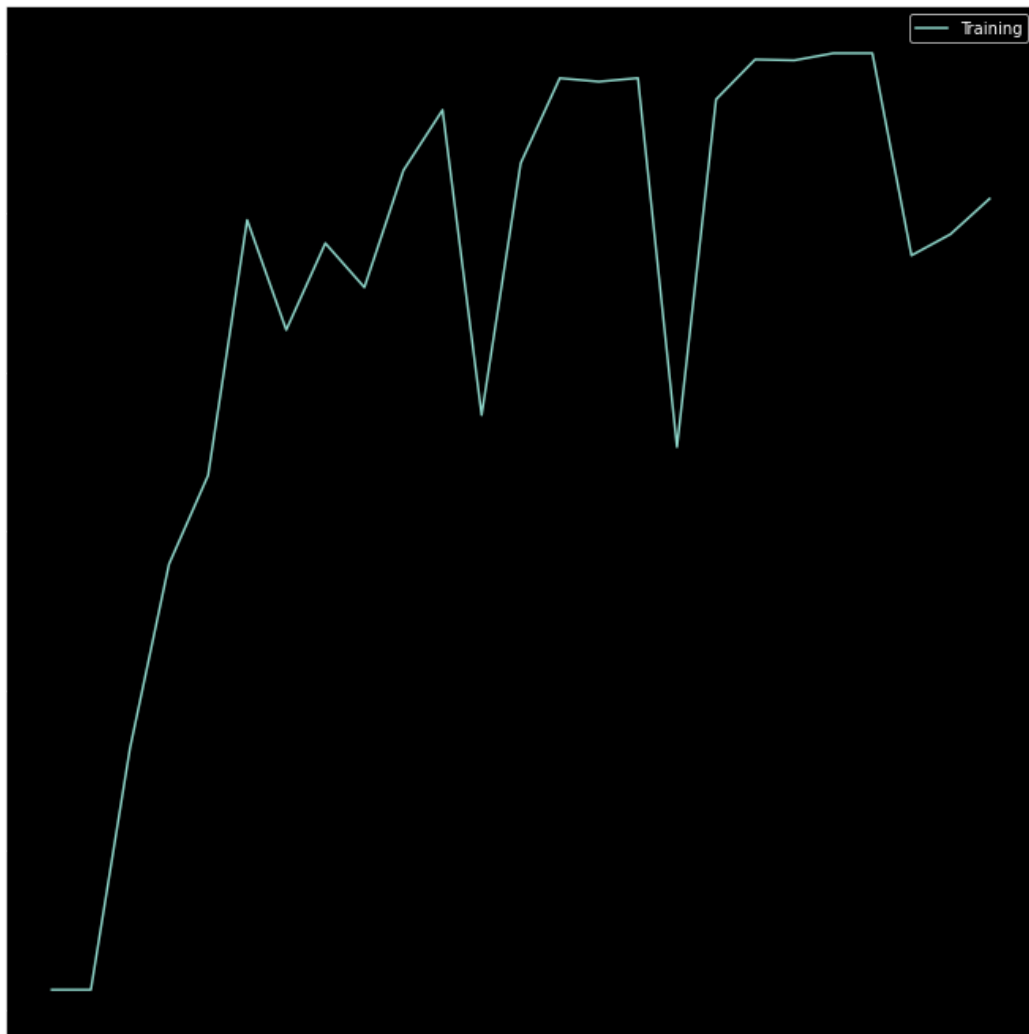
Saved Model to Disk

```
[26]: # predicting images
      path = 'F:/Courses/COURSEERA/Tensorflow in practice/Course1/Week3/human.jpg'
      img = image.load_img(path, target_size=(300, 300))
      x = image.img_to_array(img)
      x = np.expand_dims(x, axis=0)
      images = np.vstack([x])
      classes = model.predict(images, batch_size=10)
      print(classes[0])
      if classes[0]>0.5:
          print("It is a human")
      else:
          print("It is a horse")
```


[1.]

It is a human

```
[27]: plt.figure(figsize=(10,10))
      plt.style.use('dark_background')
      plt.plot(history.history['accuracy'])
      # plt.plot(history.history['val_accuracy'])
      plt.title('Model Accuracy')
      plt.ylabel('Accuracy')
      plt.xlabel('Epoch')
      plt.legend(['Training', 'Testing'])
      plt.tight_layout()
      plt.show()
```



```
[28]: plt.figure(figsize=(10,10))
plt.style.use('dark_background')
plt.plot(history.history['loss'])
# plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Training', 'Testing'])
plt.show()
```

[]: